

VigenereSolver-ng~ How it works

VigenereSolver-ng

VigenereSolver-ng is an advanced toolkit for analyzing and breaking Vigenère ciphers, designed for cryptanalysis research and educational purposes. It combines classical statistical attacks with modern, language-model-based techniques to robustly estimate key length and recover the key, even on challenging ciphertexts.

Table of Contents

- VigenereSolver-ng
 - Table of Contents
 - Key features and novel techniques
 - How to use
 - Solve ciphertexts (the usual thing)
 - Generate test data
 - Encrypt a plaintext file
 - Practical tuning knobs
 - Troubleshooting quick refs
 - 1) Model of the cipher and the text
 - 2) Coincidence, the periodogram, and key-length detection
 - 2.1 Coincidence probability at lag ℓ
 - 2.2 Windowed periodogram and variance reduction
 - 2.3 Peakiness and stability objective for w_t
 - 3) Friedman estimate and candidate key lengths
 - 4) Initial key by coset correlation (per m)
 - 5) Language model and the decryption objective

- 5.1 Interpolated Kneser–Ney (KN) probabilities
 - 5.2 Per-character negative log-likelihood (NLL)
 - 5.3 Legacy fitness and blended score
 - 6) Coordinate-wise optimization over the key
 - 6.1 Decomposition of local influence
 - 6.2 Monotone descent and convergence
 - 7) Effectiveness guarantees
 - 7.1 Consistency of key-length detection
 - 7.2 Correctness of coset shifts (initial key)
 - 7.3 Optimality of the LM objective at the true key
 - 7.4 Convergence of the key optimizer
 - 8) Auto-tuned window/step with LM tie-break
 - 9) Readability segmentation (post-processing)
 - 10) Summary of guarantees
-

Key features and novel techniques

- **Language Model Integration:** Uses high-order n-gram language models (up to 5-gram) for scoring candidate plaintexts and keys, providing much greater accuracy than traditional frequency analysis.
- **Windowed Coincidence Periodogram:** Introduces a windowed version of the coincidence periodogram, which computes coincidence rates over sliding windows to localize and stabilize key-length signals, especially on heterogeneous or short texts.
- **Key-Length Voting and Non-Maximum Suppression:** Implements a voting mechanism across windows and applies non-maximum suppression to robustly select likely key lengths, reducing false positives from harmonics and noise.
- **Jensen-Shannon Divergence Scoring:** Uses JS divergence between observed and English letter distributions to weight key character votes, improving key recovery in the presence of uneven letter frequencies.
- **Kasiski Examination with Factor Analysis:** Augments classical Kasiski examination by aggregating factors of repeated-sequence spacings, then ranks candidate key lengths by their frequency as divisors.

- **Plaintext Generator for Testing:** Includes a generator for English-like plaintexts using real language data, enabling realistic benchmarking of attacks.

These innovations make VigenereSolver-ng more effective and reliable than standard Vigenère solvers, especially on real-world ciphertexts with non-uniform content or formatting.

How to use

Requires Python 3.11+ and the project's `requirements.txt`.

Solve ciphertexts (the usual thing)

1. Install & activate env

```
python -m venv .venv
source .venv/bin/activate # Windows: .venv\Scripts\activate
pip install -r requirements.txt
```

2. Put your ciphertext in a file using triple quotes (you can have multiple blocks):

```
"""
PXWZB ... ZQL
"""

"""
ANOTHER CIPHERTEXT BLOCK ...
"""
```

3. Run the solver

```
python SolverSite/solver.py --input ciphertexts/tests.txt --passes 6 --decoder lm
```

- The solver auto-tunes the window/step, tests several key lengths in parallel, and prints the best key with IoC and score.
- Output includes a **readable plaintext** (with optional word segmentation when the original had no spaces).

Speed tip: Add `--workers <N>` to control parallelism across candidate key lengths (defaults to CPU count).

Determinism tip: For reproducible generation experiments, use `--seed <int>`; solving itself is mostly deterministic aside from small randomization in sweeps.

Generate test data

Create realistic test ciphertexts (plus sidecar keys JSON):

```
python SolverSite/solver.py --generate 5 --words 200 --out generated_ciphertexts.txt
# → ciphertexts in triple-quoted blocks
# → keys in generated_ciphertexts.txt.keys.json
```

Tweak key-length range with `--min-key` / `--max-key`.

Encrypt a plaintext file

Turn a plaintext into Vigenère ciphertext while **preserving original layout** (spacing, punctuation, case):

```
python SolverSite/solver.py --encrypt-file raw_text/kafka.txt --key SECRET --out ciphertexts/kafka_ct.txt
# If --key is omitted, a random key length [3..50] is chosen.
```

Practical tuning knobs

- **Decoder:** `--decoder lm` (default, KN 3–5-gram LM) or `--decoder legacy` (χ^2 /JSD/ngram blend).
 - **Optimization budget:** `--passes 4..8` — more passes = more key refinement (diminishing returns beyond ~6–8).
 - **Auto window/step:** on by default; to **fix** them:
`--no-auto-ws --window 600 --step 150`
 - **Annealing (escape plateaus):** `--anneal 0.05` (small positive) occasionally accepts worse local moves to avoid shallow minima.
 - **LM blend weight:** `--lm-weight 0.65` mixes LM NLL with legacy fitness for final ranking.
 - **Segmentation off:** if you prefer raw, unsegmented output: `--no-seg`.
-

Troubleshooting quick refs

- **“No ciphertext found”** → Ensure blocks are wrapped in `"" ... ""` or pass raw text as a single block.
 - **Slow on huge inputs** → Lower `--passes`, use `--workers`, or temporarily `--no-auto-ws`.
 - **Weird characters** → Save files as UTF-8. Only A–Z are analyzed; other chars are preserved in-place when formatting.
-

1) Model of the cipher and the text

Let the plaintext be a sequence of random variables P_1, P_2, \dots, P_N taking values in $\mathcal{A} = \{0, \dots, 25\}$ (A=0,...,Z=25).

A fixed Vigenère key $K = (K_0, \dots, K_{m-1}) \in \mathcal{A}^m$ produces ciphertext

$$C_i \equiv P_i + K_{i \bmod m} \pmod{26} \quad (1 \leq i \leq N).$$

We assume the letter process $\{P_i\}$ is stationary and ergodic with marginal $p(a) = \mathbb{P}\{P_i = a\}$ and higher-order statistics captured by an n -gram distribution (language model).

All non-alphabetic characters of the original message are carried through as fixed tokens; only the projected A–Z stream is transformed algebraically. This projection/reinjection step is a bijection between \mathcal{A}^N and the subset of strings with the same tokenization metadata, hence layout is exactly preserved.

2) Coincidence, the periodogram, and key-length detection

2.1 Coincidence probability at lag ℓ

Define the coincidence rate at lag ℓ :

$$R_\ell \triangleq \frac{1}{N-\ell} \sum_{i=1}^{N-\ell} \mathbf{1}\{C_i = C_{i+\ell}\}.$$

Under (1),

$$C_i = C_{i+\ell} \iff P_i + K_{i \bmod m} \equiv P_{i+\ell} + K_{(i+\ell) \bmod m} \pmod{26}.$$

- If $\ell \equiv 0 \pmod{m}$, the key offsets cancel, so

$$\mathbb{P}\{C_i = C_{i+\ell}\} = \mathbb{P}\{P_i = P_{i+\ell}\}.$$

Under an i.i.d. approximation $P_i \sim p$, the RHS equals the **index of coincidence**

$$I_c(p) \triangleq \sum_{a \in \mathcal{A}} p(a)^2.$$

For English, $I_c(p) \approx 0.066$.

- If $\ell \not\equiv 0 \pmod{m}$, let $d \equiv K_{i \bmod m} - K_{(i+\ell) \bmod m} \not\equiv 0$. Then

$$C_i = C_{i+\ell} \iff P_{i+\ell} \equiv P_i + d \pmod{26},$$

so (under i.i.d.)

$$\mathbb{P}\{C_i = C_{i+\ell}\} = \sum_a p(a) p(a + d) \triangleq S_d(p).$$

For typical language p , $S_d(p)$ is close to the random baseline $1/26 \approx 0.0385$ and strictly less than $I_c(p)$ unless p is uniform or pathologically symmetric.

Separation. For every ℓ with $\ell \equiv 0 \pmod{m}$,

$$\mathbb{E}[R_\ell] = I_c(p), \quad \text{whereas for } \ell \not\equiv 0 \pmod{m}, \quad \mathbb{E}[R_\ell] = S_d(p) < I_c(p).$$

By the strong law of large numbers (SLLN), $R_\ell \rightarrow \mathbb{E}[R_\ell]$ almost surely as $N \rightarrow \infty$. Hence the periodogram $\ell \mapsto R_\ell$ exhibits peaks at $\ell \in m\mathbb{Z}_+$.

2.2 Windowed periodogram and variance reduction

Let $\{(s_j, s_j + w - 1)\}$ be sliding windows of length w and stride t . Define the windowed coincidence:

$$R_\ell^{(j)} \triangleq \frac{1}{w - \ell} \sum_{i=s_j}^{s_j+w-\ell-1} \mathbf{1}\{C_i = C_{i+\ell}\}.$$

The **window-averaged** periodogram is

$$\overline{R}_\ell \triangleq \frac{1}{J} \sum_{j=1}^J R_\ell^{(j)}.$$

Assuming stationarity within windows, $\mathbb{E}[\overline{R}_\ell] = \mathbb{E}[R_\ell]$, and

$$\text{Var}(\overline{R}_\ell) = \frac{1}{J^2} \sum_{j=1}^J \text{Var}(R_\ell^{(j)}) + \frac{2}{J^2} \sum_{j < k} \text{Cov}(R_\ell^{(j)}, R_\ell^{(k)}).$$

For modest overlap (stride t not too small), the covariance terms are bounded and $\text{Var}(\overline{R}_\ell) = O(1/J)$. Thus window averaging lowers estimator variance and stabilizes peak detection under topic/style drift.

2.3 Peakiness and stability objective for w, t

Let $\Phi(\bar{R})$ denote a **peakiness statistic**, e.g.

$$\Phi(\bar{R}) \triangleq \frac{\frac{1}{k} \sum_{r=1}^k \bar{R}_{\ell_{(r)}} - \text{median}_{\ell} \bar{R}_{\ell}}{\sqrt{\text{Var}_{\ell}(\bar{R}_{\ell}) + \varepsilon}},$$

where $\ell_{(1)}, \dots, \ell_{(k)}$ are the top k lags.

Partition the ciphertext into thirds and compute the top- k lag sets T_1, T_2, T_3 . Define stability

$$\Psi \triangleq \frac{1}{3} \sum_{1 \leq u < v \leq 3} \frac{|T_u \cap T_v|}{|T_u \cup T_v|}.$$

Add a weak prior Υ encouraging $\arg \max_{\ell} \bar{R}_{\ell}$ to be near a multiple of the Friedman guess \tilde{m} , e.g.

$$\Upsilon \triangleq \frac{1}{1 + \min_{k \in \mathbb{N}} |\hat{\ell} - k\tilde{m}|}, \quad \hat{\ell} \in \arg \max_{\ell} \bar{R}_{\ell},$$

and a utility term U penalizing degenerate windows (too short/long for N). The tuning objective is

$$J(w, t) \triangleq \alpha \Phi(\bar{R}) + \beta \Psi + \gamma \Upsilon + \eta U, \quad \alpha, \beta, \gamma, \eta > 0.$$

Maximizing $J(w, t)$ increases the signal-to-noise of the key-period peaks and their reproducibility across segments. A tie-break among top (w, t) uses the language-model score defined in §4 to prefer pairs that make subsequent decryption “look” more like the target language.

3) Friedman estimate and candidate key lengths

Let $I_c(\text{obs})$ be the observed loC of the cleaned ciphertext and let $I_r = 1/26$ be the random baseline. With $I_e = I_c(p)$ the English loC, the (classical) Friedman estimate is

$$\hat{m}_{\text{Friedman}} \approx \frac{I_e - I_r}{I_c(\text{obs}) - I_r}.$$

Rounding and taking the neighborhood $\{\widehat{m}_F - 2, \dots, \widehat{m}_F + 2\}$, then uniting with the top lags of \overline{R}_ℓ and the small-factor set from Kasiski's test yields a finite candidate set $\mathcal{M} \subset \{2, \dots, 50\}$.

By §2.1–2.2 and LLN, if N is large enough, the true m appears among the top lags with probability $\rightarrow 1$, and thus in \mathcal{M} .

4) Initial key by coset correlation (per m)

Fix $m \in \mathcal{M}$. Split the cleaned ciphertext into cosets

$$C^{(r)} \triangleq (C_r, C_{r+m}, C_{r+2m}, \dots), \quad r = 0, \dots, m-1.$$

Within a sliding window W of length w , let $\hat{f}^{(r)} \in \Delta^{25}$ be the empirical histogram of $C^{(r)}$ letters (as \mathcal{A} values). For a Caesar shift $s \in \mathcal{A}$, define the shifted histogram $\hat{f}_{-s}^{(r)}$ by $(\hat{f}_{-s}^{(r)})_a = \hat{f}_{a+s}^{(r)}$. With a reference English distribution $q \in \Delta^{25}$, consider

$$s^* \in \arg \max_{s \in \mathcal{A}} \langle \hat{f}_{-s}^{(r)}, q \rangle = \arg \min_{s \in \mathcal{A}} H(\hat{f}_{-s}^{(r)}, q),$$

where $H(p, q) = -\sum_a p(a) \log q(a)$ is cross-entropy.

Claim (MLE for a Caesar coset). If the plaintext letters in coset r are i.i.d. $\sim q$, then s^* maximizes the log-likelihood of the observed coset under a Caesar model.

Proof. The log-likelihood for shift s is $\sum_a n_a^{(r)} \log q(a - s)$, which equals $n \langle \hat{f}_{-s}^{(r)}, \log q \rangle$. Maximizing this is equivalent to minimizing $H(\hat{f}_{-s}^{(r)}, q)$. \square

Windows whose letter distribution deviates from English are down-weighted using a divergence D (here Jensen–Shannon): with weights $w_W \propto 1/D(\hat{f}_W, q)$, the final per-coset shift is the weighted mode over windows. This reduces variance by emphasizing windows closer to the stationary regime.

5) Language model and the decryption objective

5.1 Interpolated Kneser–Ney (KN) probabilities

Let $c_n(\cdot)$ be counts for n -grams over \mathcal{A} , and $\text{cont}_n(h)$ the number of unique continuations of a context h (order $n - 1$). With absolute discount $D \in (0, 1)$, the KN conditional probability for a next symbol w given context h of length $n - 1$ is

$$p_{\text{KN}}(w \mid h) = \frac{\max\{c_n(hw) - D, 0\}}{c_{n-1}(h\cdot)} + \lambda(h) p_{\text{KN}}(w \mid h'), \quad \lambda(h) \triangleq \frac{D \text{cont}_n(h)}{c_{n-1}(h\cdot)},$$

with backoff to h' the suffix of h . For the base case $n = 1$, the continuation probability can be taken as

$$p_{\text{cont}}(w) \triangleq \frac{N_{1+}(*, w)}{N_{1+}(*, *)}, \quad N_{1+}(*, w) = |\{x \in \mathcal{A} : c_2(xw) > 0\}|, \quad N_{1+}(*, *) = |\{(x, y) \in \mathcal{A}^2 : c_2(xy) > 0\}|.$$

5.2 Per-character negative log-likelihood (NLL)

For a cleaned candidate plaintext $x_1^N \in \mathcal{A}^N$ and maximum order $n_{\max} = 5$,

$$\mathcal{L}(x_1^N) \triangleq \frac{1}{N} \sum_{i=1}^N -\log p_{\text{KN}}(x_i \mid x_{i-k}^{i-1}), \quad k = \min\{n_{\max} - 1, i - 1\}.$$

Given a key $K \in \mathcal{A}^m$, define the decryption mapping $\text{Dec}_K(C)_i \equiv C_i - K_{i \bmod m} \pmod{26}$ and the **LM objective**

$$F(K) \triangleq \mathcal{L}(\text{Dec}_K(C)).$$

5.3 Legacy fitness and blended score

Let $\text{fit}(x_1^N)$ be a convex combination of reduced χ^2 , Jensen–Shannon divergence, and n -gram surprisals (3–4 grams) on x_1^N . The **selection score** is

$$S(K) \triangleq w F(K) + (1 - w) \text{fit}(\text{Dec}_K(C)), \quad w \in (0, 1].$$

6) Coordinate-wise optimization over the key

6.1 Decomposition of local influence

Let $m = |K|$. For a residue class (coset) $r \in \{0, \dots, m-1\}$, write $\mathcal{I}_r = \{i : i \equiv r \pmod{m}\}$. Consider modifying only K_r to $s \in \mathcal{A}$. Let $x^{(s)}$ be the plaintext after this modification. For a fixed n_{\max} , the change in F obeys

$$F(K_{[r \leftarrow s]}) - F(K) = \frac{1}{N} \sum_{i \in \mathcal{N}_r} \left[-\log p_{\text{KN}}(x_i^{(s)} \mid x_{i-k}^{(s) \ i-1}) + \log p_{\text{KN}}(x_i \mid x_{i-k}^{i-1}) \right],$$

where \mathcal{N}_r is the set of indices whose n_{\max} -gram window **touches** any position in \mathcal{I}_r . Thus only those local windows need to be rescored when optimizing K_r .

6.2 Monotone descent and convergence

Define one **coordinate update** at residue r by

$$K_r \leftarrow \arg \min_{s \in \mathcal{A}} F(K_{[r \leftarrow s]}).$$

(This is the **anneal=0** case.) Since the minimizer is selected exactly over a finite set,

$$F(K^{(t+1)}) \leq F(K^{(t)}) \quad \text{at every step,}$$

hence $\{F(K^{(t)})\}$ is a bounded, monotonically non-increasing sequence and converges. Because there are only 26^m keys, and strict decreases can occur only finitely many times, the process terminates at a **coordinate-wise minimum** K^* (a fixed point of all residue-wise updates).

When a small simulated annealing acceptance is used, monotonicity is relaxed to escape plateaus; nonetheless the state space is finite, and with a standard cooling schedule the process converges almost surely to a local minimum.

7) Effectiveness guarantees

We state assumptions explicitly:

- **A1 (Language).** The true plaintext P_1^N is generated by a stationary ergodic process whose n -gram statistics are well-approximated by the KN model used to score text.
- **A2 (Key).** The Vigenère key $K \in \mathcal{A}^m$ is fixed and unknown; m is bounded (e.g., ≤ 50).
- **A3 (Non-degenerate alphabets).** The cleaned text length $N \rightarrow \infty$ and letter frequency vector is non-uniform (true for English).

7.1 Consistency of key-length detection

Theorem 1 (Periodogram consistency). Under A1–A3, the set of top lags of \overline{R}_ℓ contains m and its multiples with probability $\rightarrow 1$ as $N \rightarrow \infty$. Consequently, the candidate pool \mathcal{M} includes m with probability $\rightarrow 1$.

Sketch. §2.1 showed $\mathbb{E}[\overline{R}_\ell] = I_c(p)$ for $\ell \in m\mathbb{Z}_+$ and $S_d(p)$ otherwise with strict separation. By SLLN, $\overline{R}_\ell \rightarrow \mathbb{E}[\overline{R}_\ell]$ uniformly over a finite set of lags. Therefore the top lags converge to the maximizers, which include multiples of m . \square

7.2 Correctness of coset shifts (initial key)

Proposition 2 (Coset MLE). Within a coset r , the shift

$$\hat{s}_r \in \arg \min_s H(\hat{f}_{-s}^{(r)}, q)$$

is the maximum-likelihood estimate of the Caesar offset assuming the coset plaintext is i.i.d. $\sim q$.

Proof. Already given in §4. \square

Variance reduction. If $D(\hat{f}_W, q)$ is a proper divergence and $\mathbb{E}[D(\hat{f}_W, q)]$ increases with departure from stationarity, then weighting windows by $w_W \propto 1/D(\hat{f}_W, q)$ yields a **minimum variance unbiased** estimator among the class of linear unbiased combinations of $\hat{s}_{r,W}$ under a heteroskedastic model of window quality. (This follows from generalized least squares: inverse-variance weights are optimal; here D is used as a monotone proxy for variance.)

7.3 Optimality of the LM objective at the true key

Let P denote the distribution of true plaintext strings, Q_K the distribution induced by decrypting with key K (wrong alignment mixes cosets and introduces Caesar-shifted contexts). Let M denote the KN model used for scoring. The per-character cross-

entropy under M is

$$H(P; M) \triangleq \lim_{N \rightarrow \infty} \mathbb{E}[\mathcal{L}(P_1^N)], \quad H(Q_K; M) \triangleq \lim_{N \rightarrow \infty} \mathbb{E}[\mathcal{L}(\tilde{P}_1^N)],$$

where $\tilde{P} = \text{Dec}_K(C)$.

Theorem 3 (Asymptotic separation). If $M = P$ (oracle LM), then for every $K \neq K_{\text{true}}$,

$$H(Q_K; M) - H(P; M) = \text{KL}(P \| Q_K) > 0.$$

Proof. By the information inequality,

$$\mathbb{E}_P[-\log P(X)] \leq \mathbb{E}_P[-\log Q_K(X)],$$

with equality iff $Q_K = P$. But $Q_K = P$ only when $K = K_{\text{true}}$ up to inconsequential Caesar offsets that are already absorbed in K ; otherwise different cosets and contexts are mismatched and the equality fails. \square

Corollary 4 (Consistency with approximate LM). If M is a consistent estimator of P in the sense that $H(P; M) \rightarrow H(P; P)$ and, uniformly over wrong keys K , $H(Q_K; M) \rightarrow H(Q_K; P)$, then there exists $\delta > 0$ such that for all large N ,

$$\mathbb{E}[F(K_{\text{true}})] + o(1) \leq \mathbb{E}[F(K)] - \delta \quad \text{for all } K \neq K_{\text{true}}.$$

Hence the LM part of the selection prefers the true key with probability $\rightarrow 1$.

Blended score. Because $S(K) = wF(K) + (1 - w)\text{fit}(\cdot)$ with $w > 0$, the separation provided by F dominates for large N , so the blended score also correctly ranks the true key with probability $\rightarrow 1$.

7.4 Convergence of the key optimizer

Proposition 5 (Finite convergence to a coordinate-wise minimum). The coordinate-wise minimization of $F(K)$ over residues $r = 0, \dots, m - 1$ with exact per-residue minimizers terminates in finitely many steps at a key K^* such that for each residue r ,

$$K_r^* \in \arg \min_{s \in \mathcal{A}} F(K_{[r \leftarrow s]}^*).$$

Proof. Each successful update strictly decreases F by at least a positive amount (finite score space due to finite data and finite alphabet), and there are finitely many keys 26^m . Hence only finitely many strict decreases occur; the process halts when no per-

residue improvement exists. \square

Locality correctness. The update uses exactly the set \mathcal{N}_r of windows affected by residue r (see §6.1). Therefore each residue update is an **exact** coordinate minimization of F , not a heuristic.

8) Auto-tuned window/step with LM tie-break

Given the objective $J(w, t)$ from §2.3, select the top few (w, t) pairs by J . For each, take the best 1–2 candidate key lengths m by \overline{R}_ℓ , form initial keys (§4), decrypt, and measure $-F(K)$ (i.e., LM likelihood). The final choice is

$$(w^*, t^*) \in \arg \max_{(w, t) \in \text{beam}} \left(\alpha' J(w, t) + \beta' \max_{m \in \mathcal{M}(w, t)} \max_{K \in \mathcal{K}_{\text{init}}(m)} -F(K) \right),$$

with $\alpha', \beta' > 0$. Under Theorem 3/Corollary 4, the LM tie-break favors (w, t) that sharpen the m proposals and produce plaintext closer to P , improving the probability that the downstream key optimization starts in the basin of attraction of K_{true} .

9) Readability segmentation (post-processing)

When the recovered clean plaintext x_1^N lacks spaces, a dictionary-rank cost $c(w) \propto \log \text{rank}(w)$ and a length penalty $\lambda|w|$ define the segmentation objective

$$\min_{x_1^N = w_1 \cdots w_T} \sum_{t=1}^T (c(w_t) + \lambda|w_t|),$$

solved by standard dynamic programming. This does not affect correctness; it only improves human readability.

10) Summary of guarantees

- **Key length:** \overline{R}_ℓ consistently peaks at multiples of m (§2), so $m \in \mathcal{M}$ w.h.p. for large N (Theorem 1).
- **Initial key:** per-coset Caesar shifts via cross-entropy minimization are MLEs (Proposition 2); divergence-weighted voting reduces variance.
- **Key optimization:** exact coordinate updates monotonically decrease F and converge to a coordinate-wise optimum (Proposition 5); only locally affected LM windows are rescored (§6.1).
- **Selection:** if the LM matches the plaintext process, the true decryption minimizes the expected NLL; any wrong key incurs a strictly larger cross-entropy (Theorem 3), and this separation persists (Corollary 4). With $w > 0$, the blended score inherits this separation.

Under these conditions, the pipeline is **asymptotically effective**: as $N \rightarrow \infty$, it recovers the correct key length with probability $\rightarrow 1$, and among candidate keys, the LM objective (and thus the blended score) is uniquely minimized at the true key.